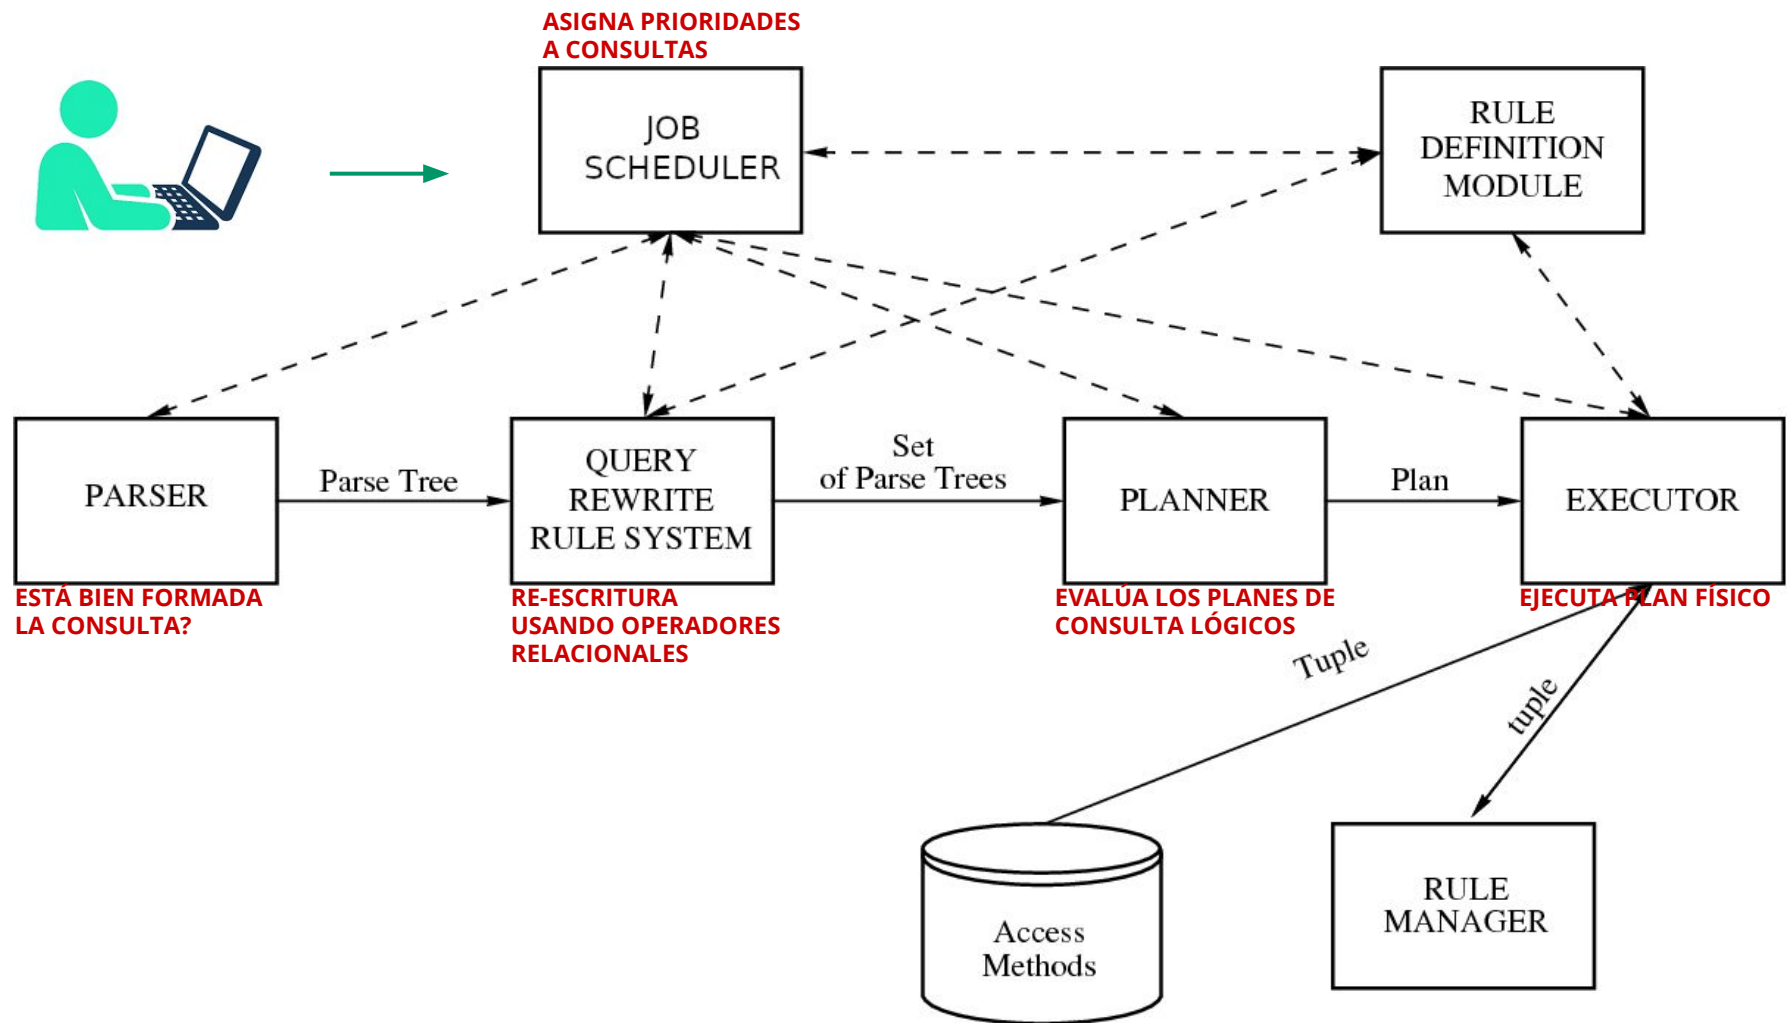
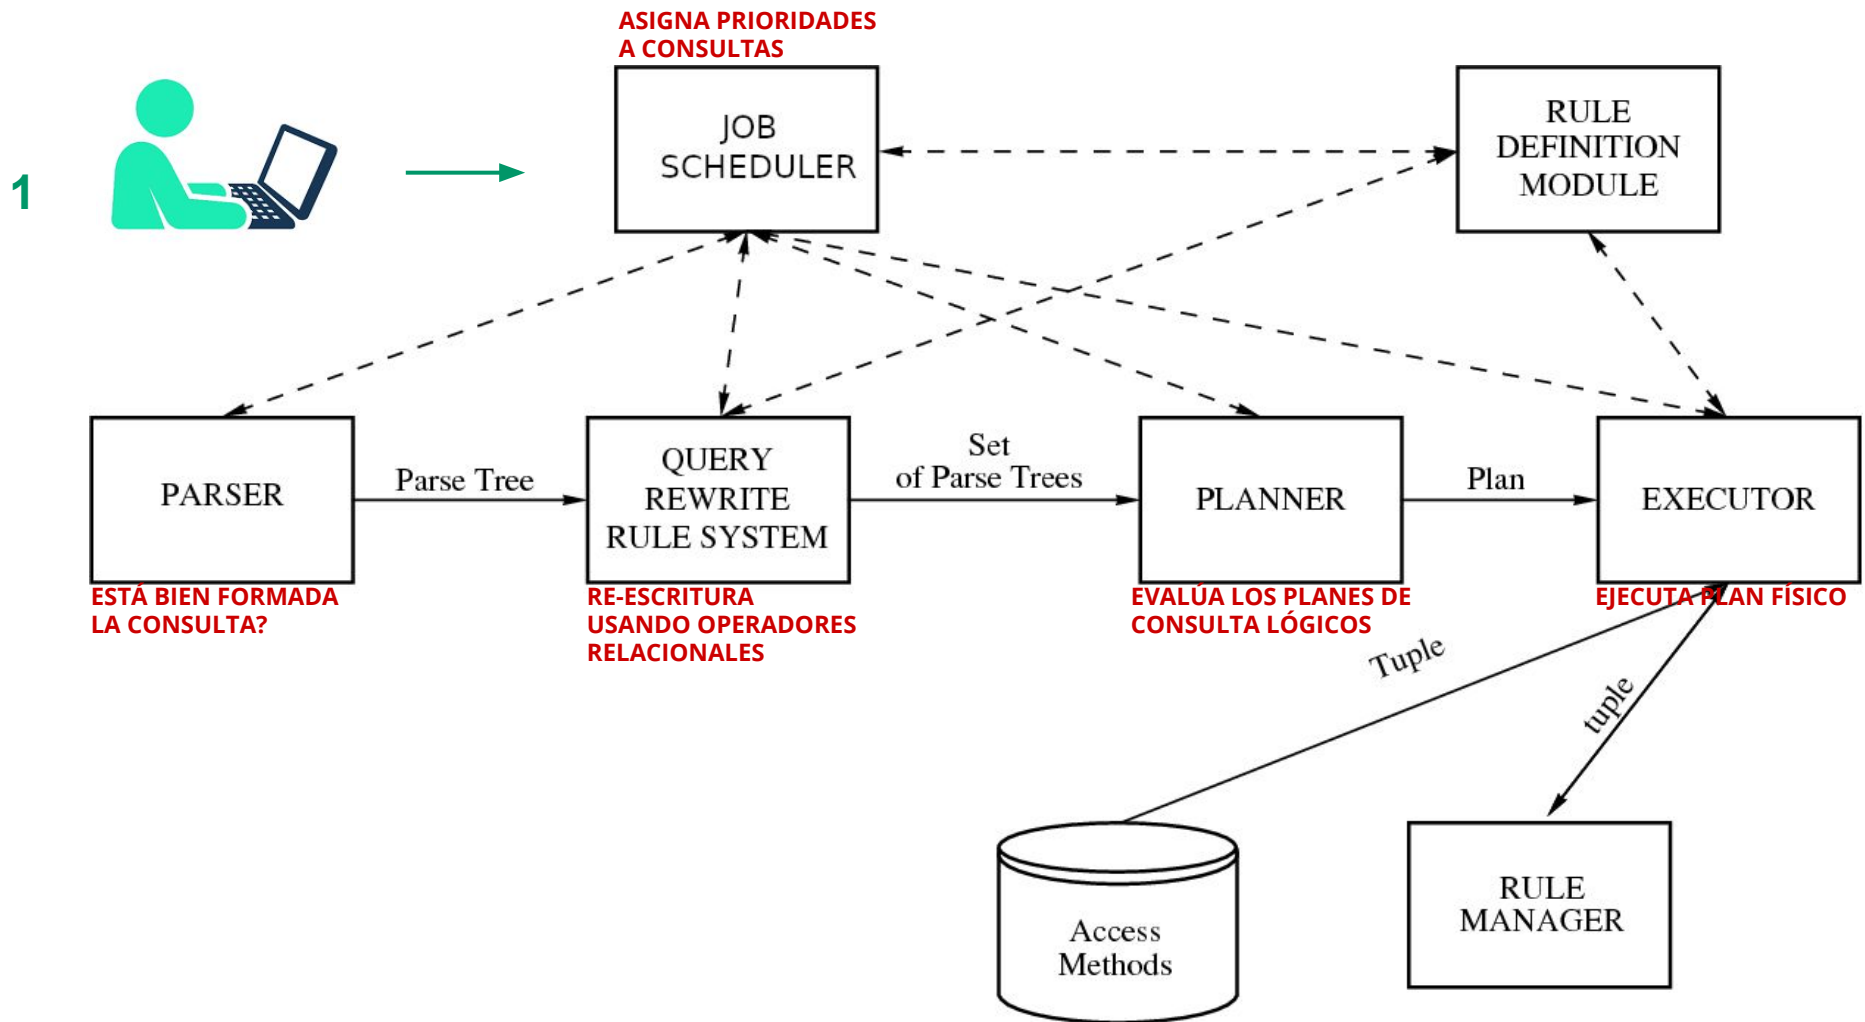

Bases de Datos

— Ejecución de consultas —

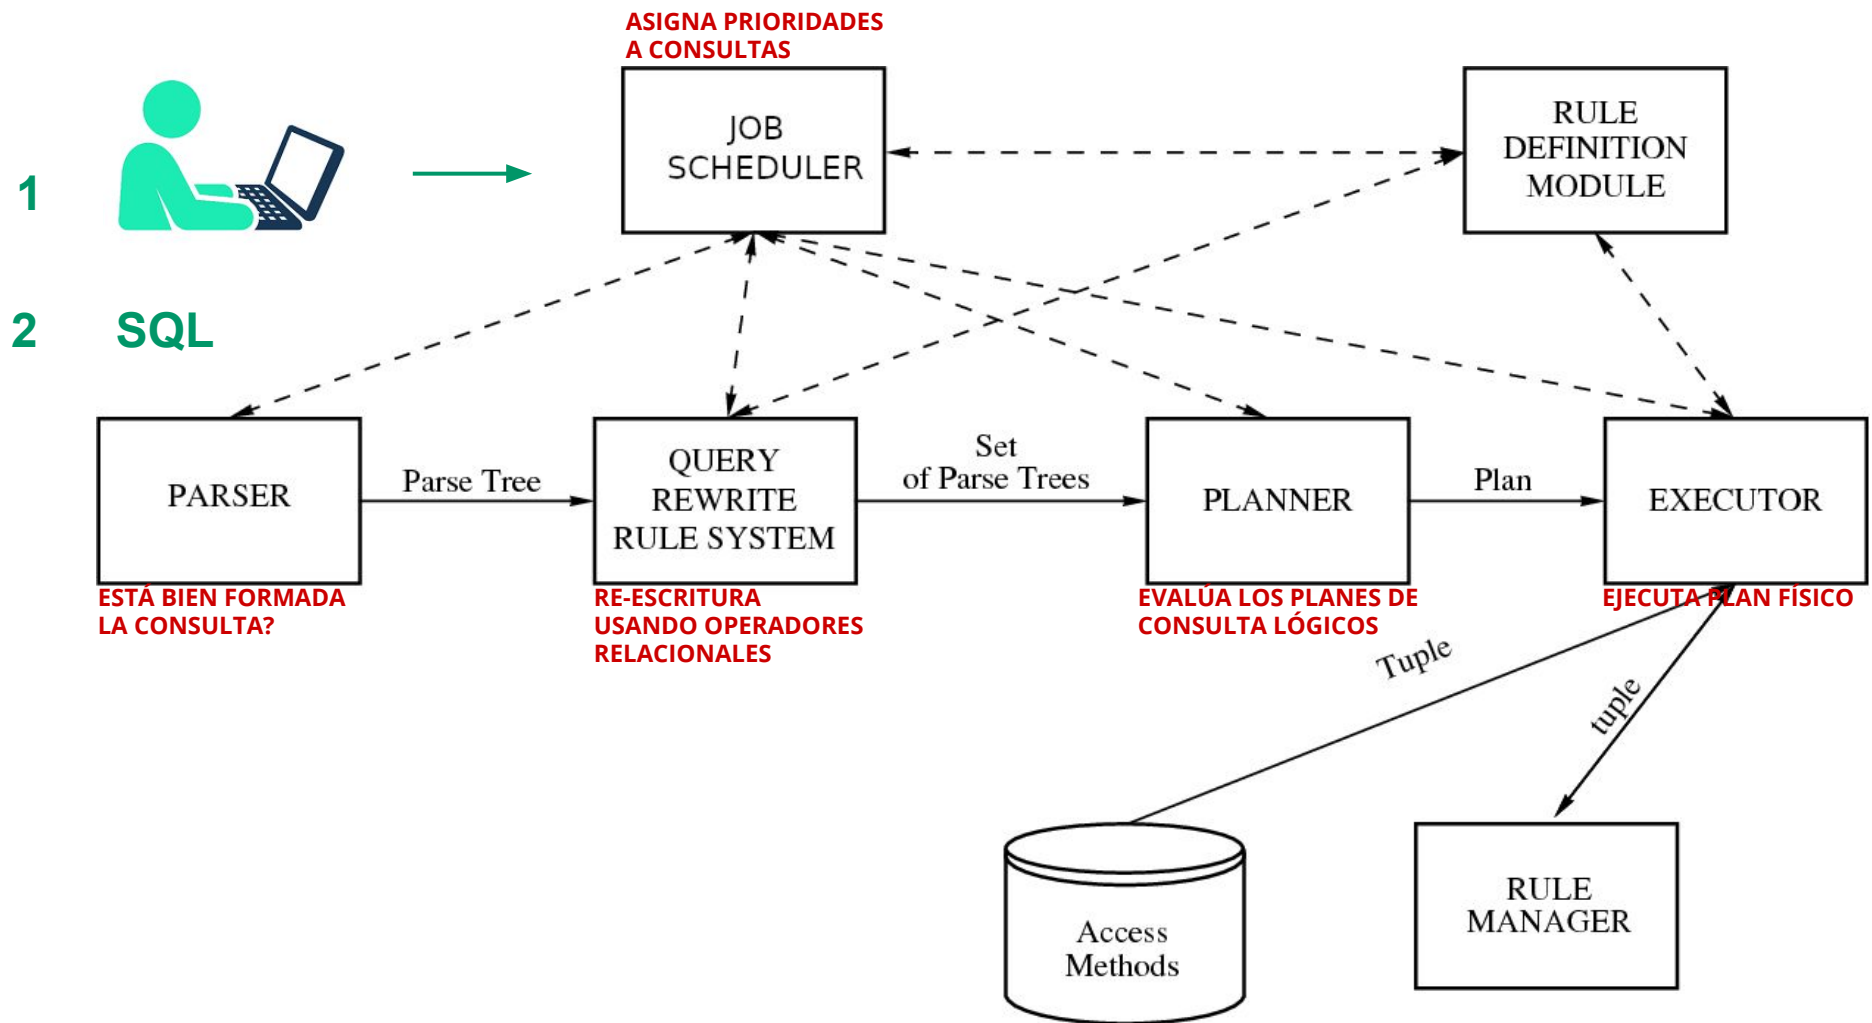
Metodología



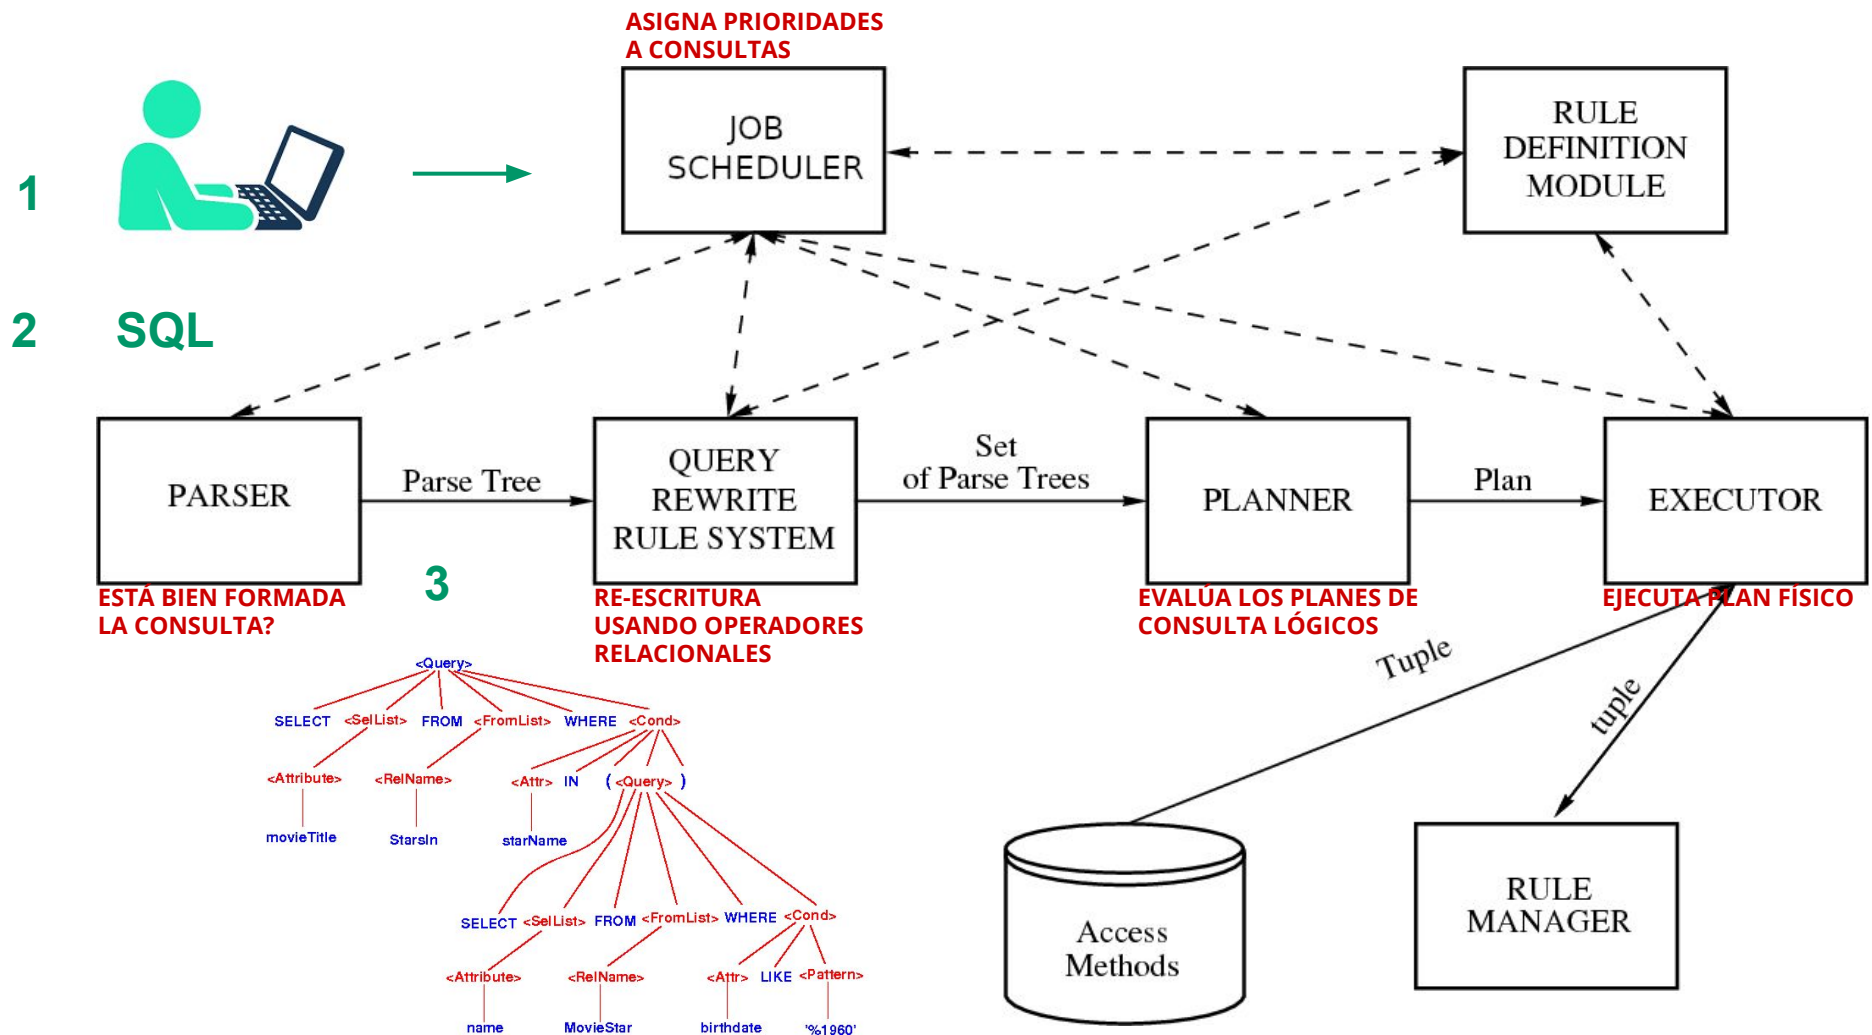
Metodología



Metodología



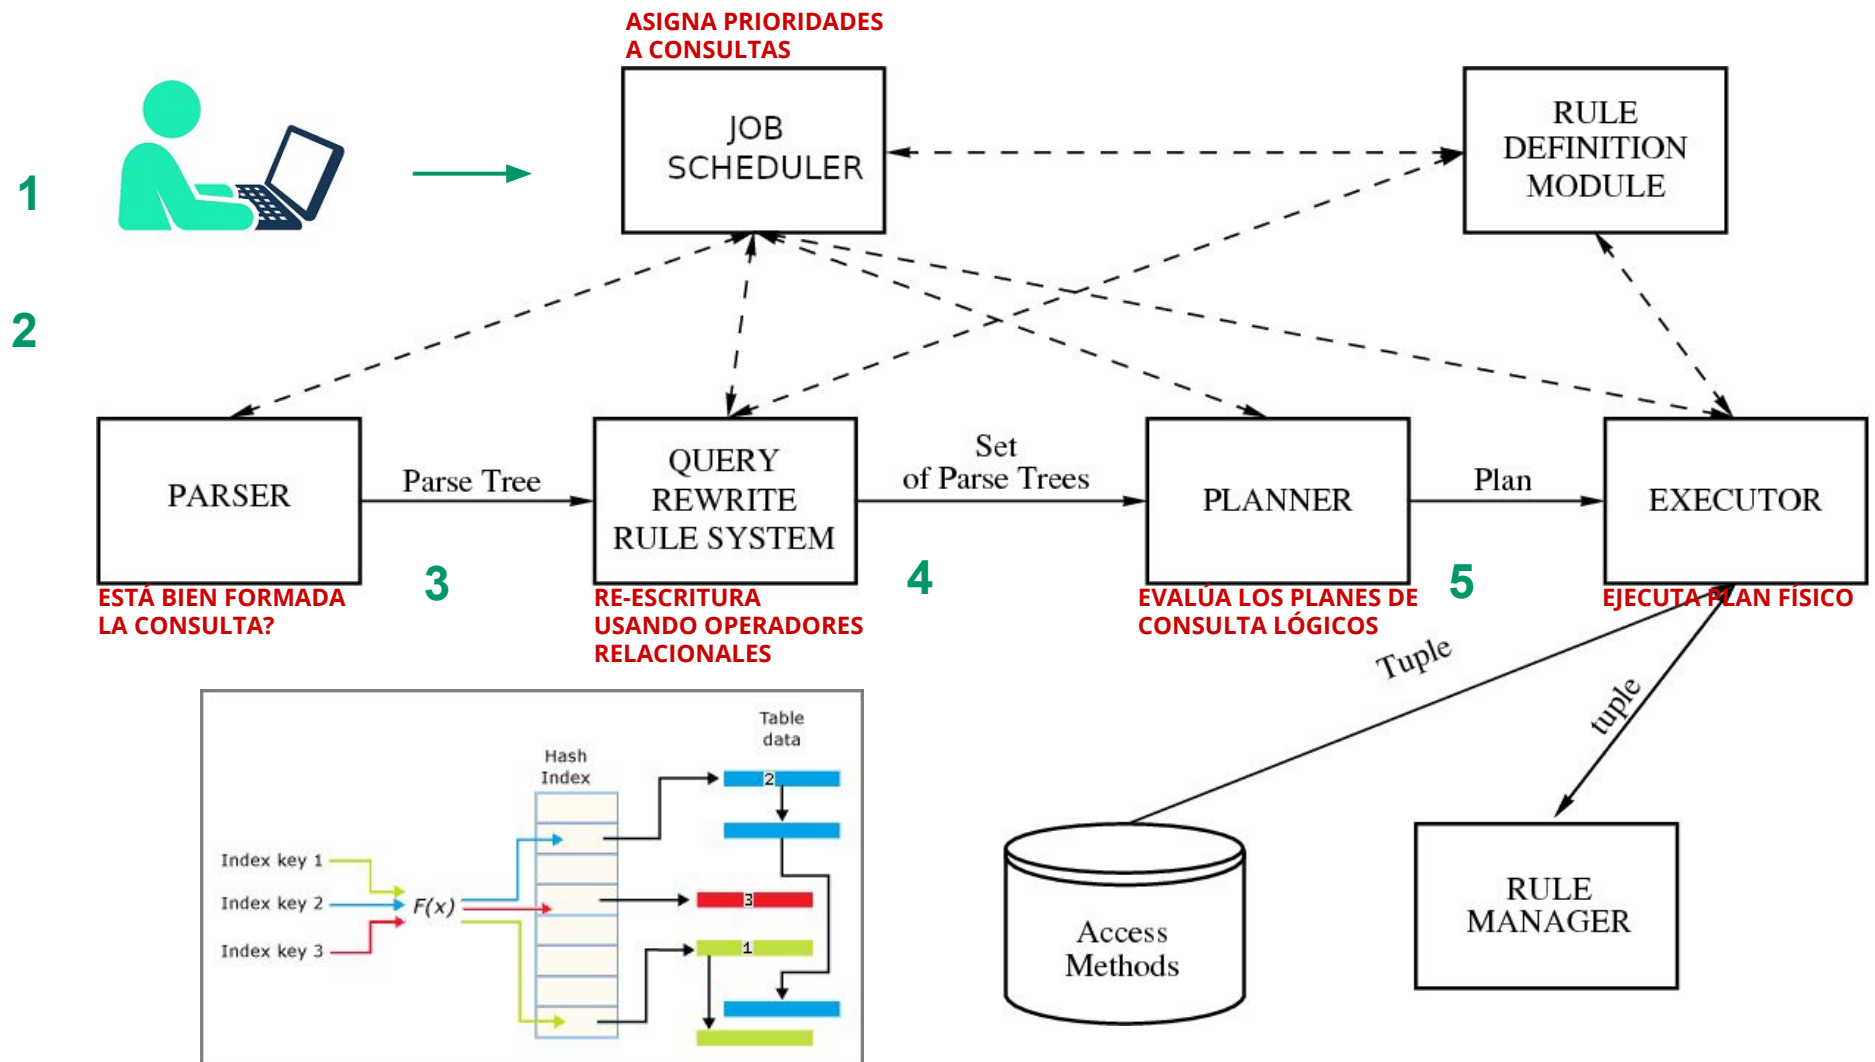
Metodología



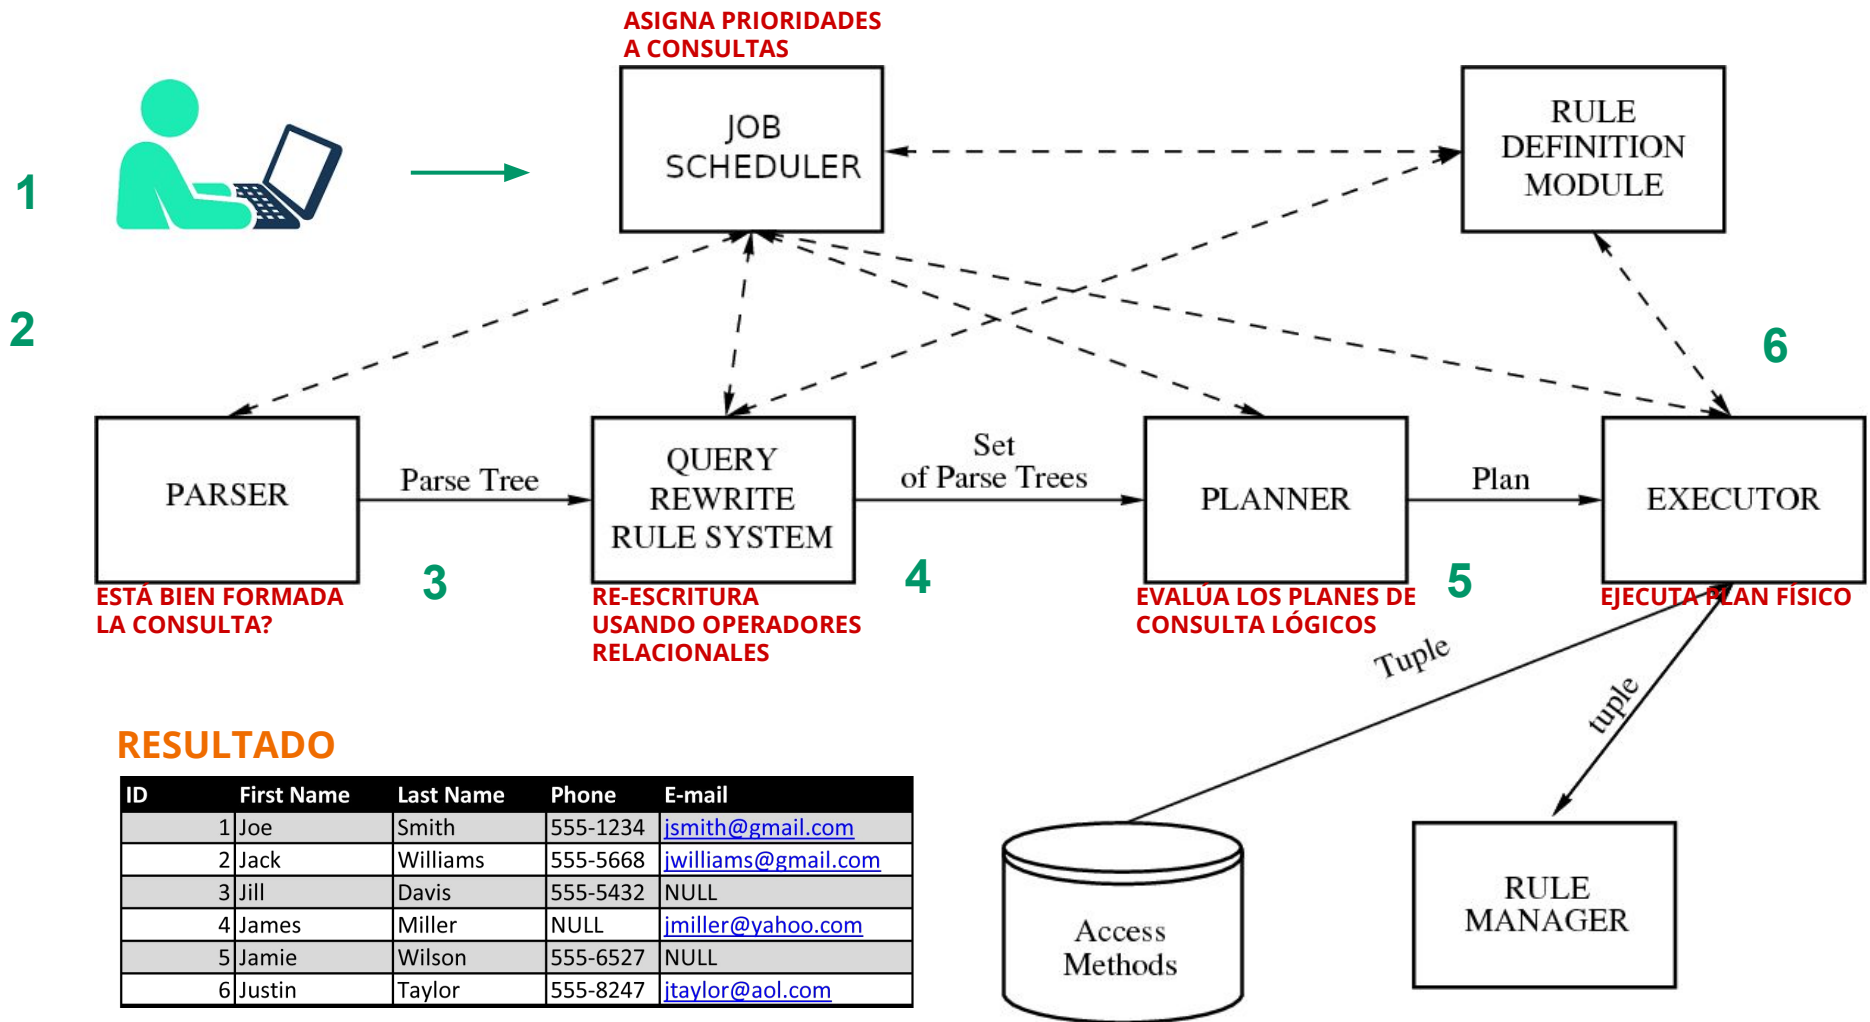
2



Metodología



Metodología



SQL Parser

Gramática SQL

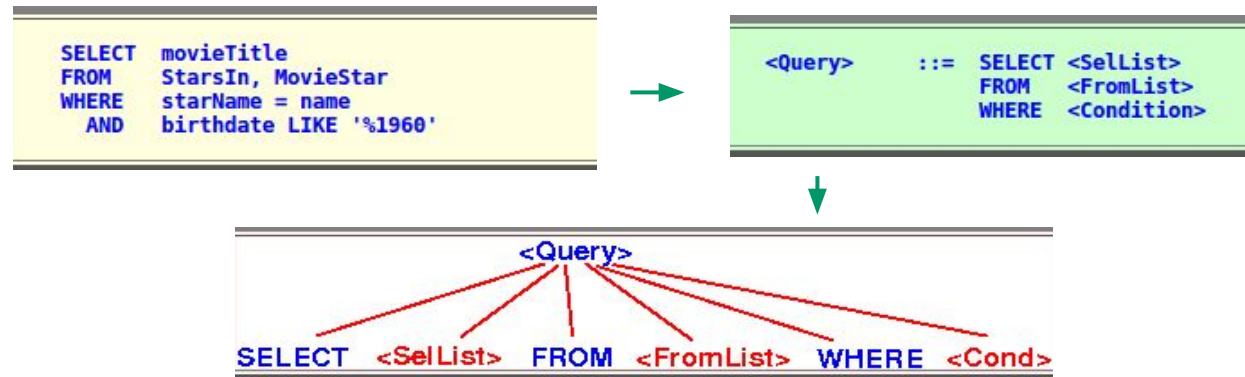
```
<Query>      ::=  SELECT <SelList>
                  FROM  <FromList>
                  WHERE  <Condition>

<SelList>     ::=  <Attribute> |
                  <Attribute> , <SelList>

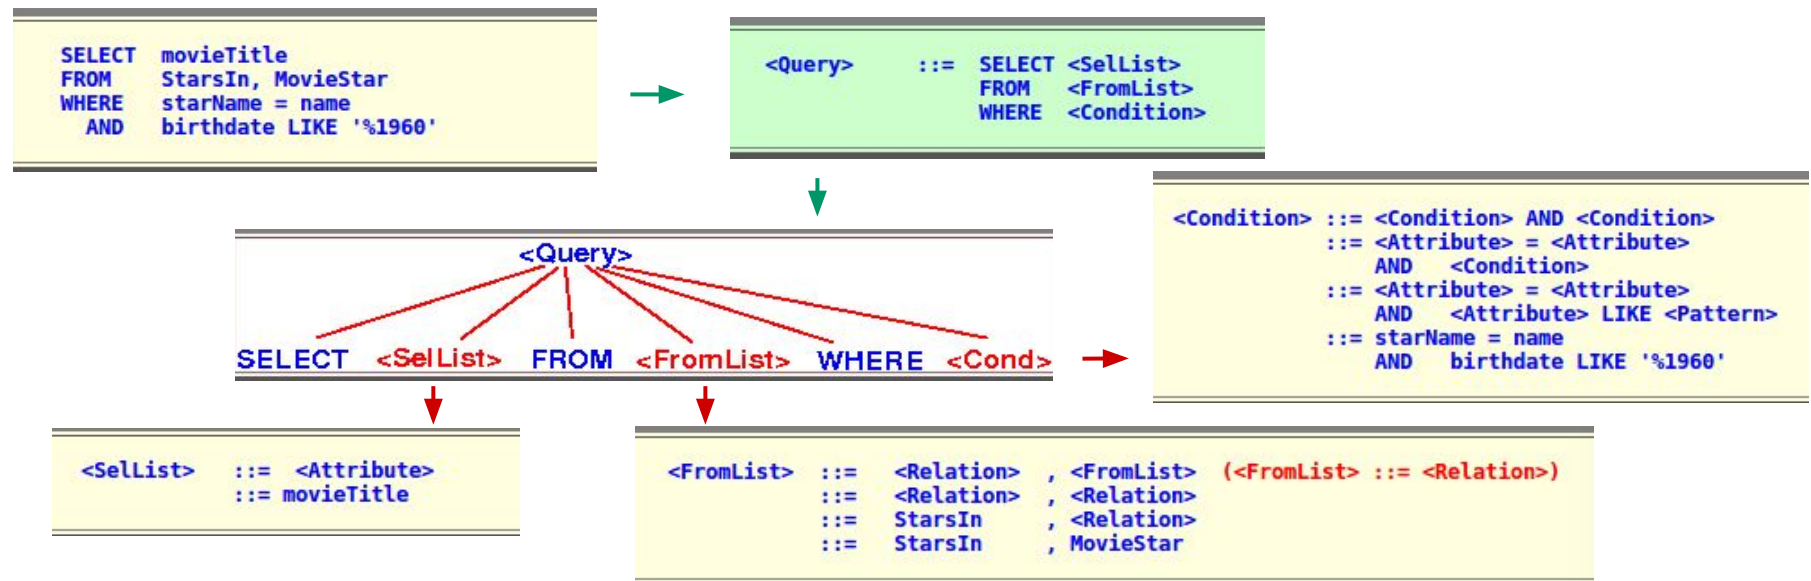
<FromList>    ::=  <Relation> |
                  <Relation> , <FromList>

<Condition>   ::=  <Condition> AND <Condition> |
                  <Attribute> IN ( <Query> ) |
                  <Attribute> = <Attribute> |
                  <Attribute> LIKE <Pattern>
```

SQL Parser



SQL Parser



SQL Parser

```
SELECT movieTitle
FROM StarsIn, MovieStar
WHERE starName = name
AND birthdate LIKE '%1960'
```

```
<Query> ::= SELECT <SelList>
          FROM  <FromList>
          WHERE <Condition>
```

```

      <Query>
    /  |  |  |  |  |
SELECT <SelList> FROM <FromList> WHERE <Cond>
```

```
<Condition> ::= <Condition> AND <Condition>
             ::= <Attribute> = <Attribute>
                AND <Condition>
             ::= <Attribute> = <Attribute>
                AND <Attribute> LIKE <Pattern>
             ::= starName = name
                AND birthdate LIKE '%1960'
```

```
<SelList> ::= <Attribute>
           ::= movieTitle
```

```
<FromList> ::= <Relation> , <FromList> (<FromList> ::= <Relation>)
           ::= <Relation> , <Relation>
           ::= StarsIn , <Relation>
           ::= StarsIn , MovieStar
```

```

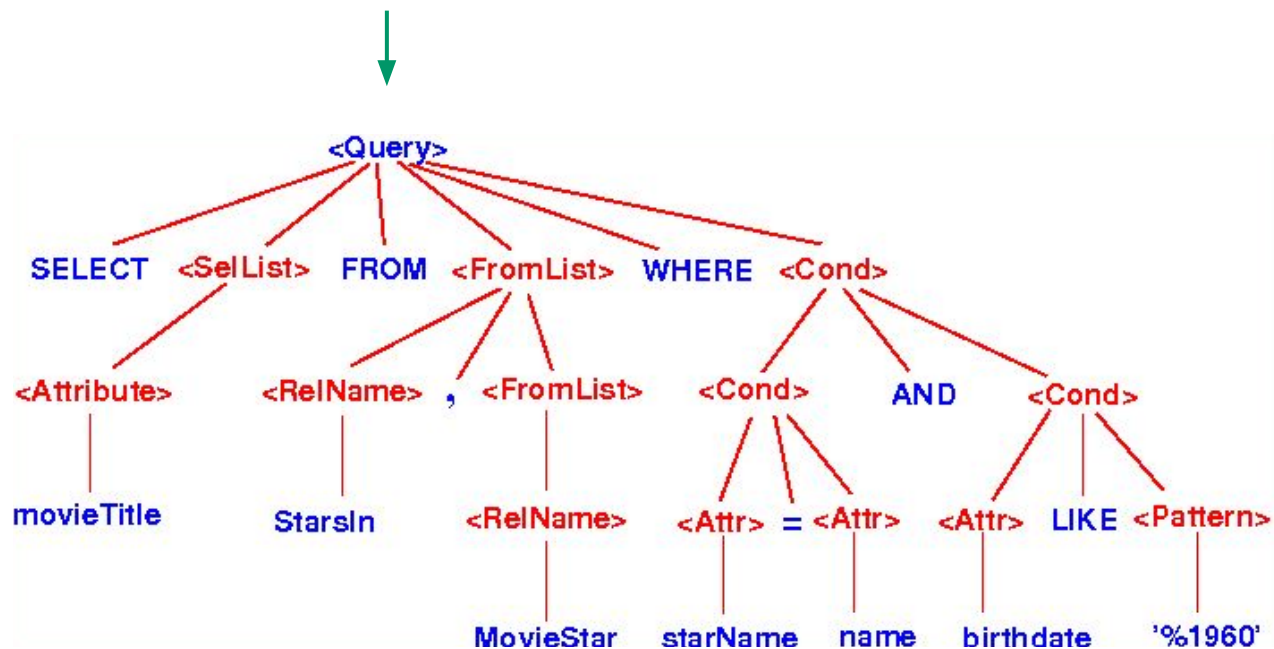
      <Query>
    /  |  |  |  |  |
SELECT <SelList> FROM <FromList> WHERE <Cond>
    /
  <Attribute>
    |
movieTitle
```

```

      <Query>
    /  |  |  |  |  |
SELECT <SelList> FROM <FromList> WHERE <Cond>
    /  |  |
  <Attribute> <RelName> , <FromList>
    |         |         |
movieTitle StarsIn <RelName>
                  |
                MovieStar
```

SQL Parser

```
SELECT movieTitle
FROM StarsIn, MovieStar
WHERE starName = name
AND birthdate LIKE '%1960'
```



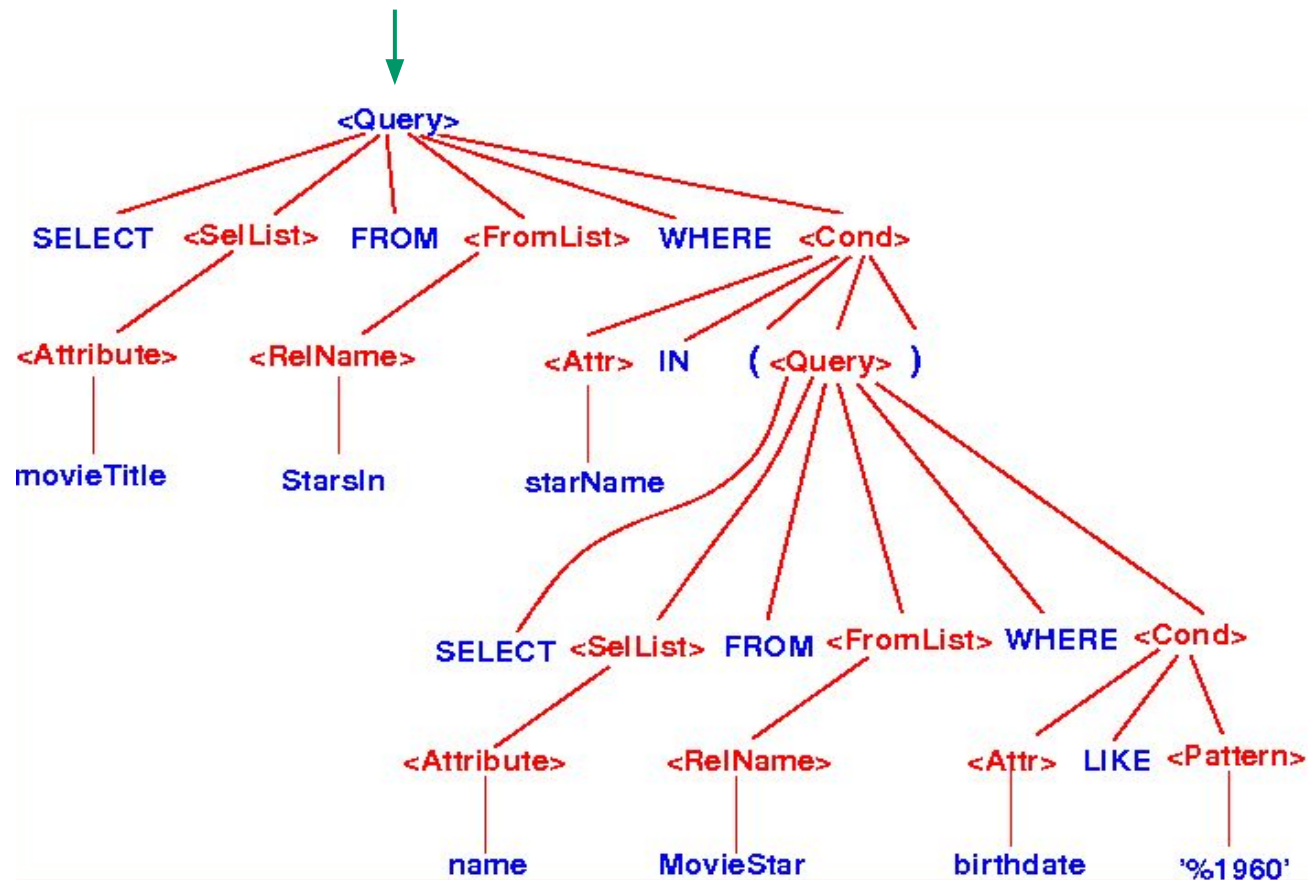
SQL Parser

```
SELECT movieTitle
FROM StarsIn
WHERE starName IN ( SELECT name
                     FROM MovieStar
                     WHERE birthdate LIKE '%1960' )
```



SQL Parser

```
SELECT movieTitle
FROM StarsIn
WHERE starName IN ( SELECT name
                     FROM MovieStar
                     WHERE birthdate LIKE '%1960' )
```



Query Rewrite

Los planes de ejecución lógicos se expresan usando operadores relacionales

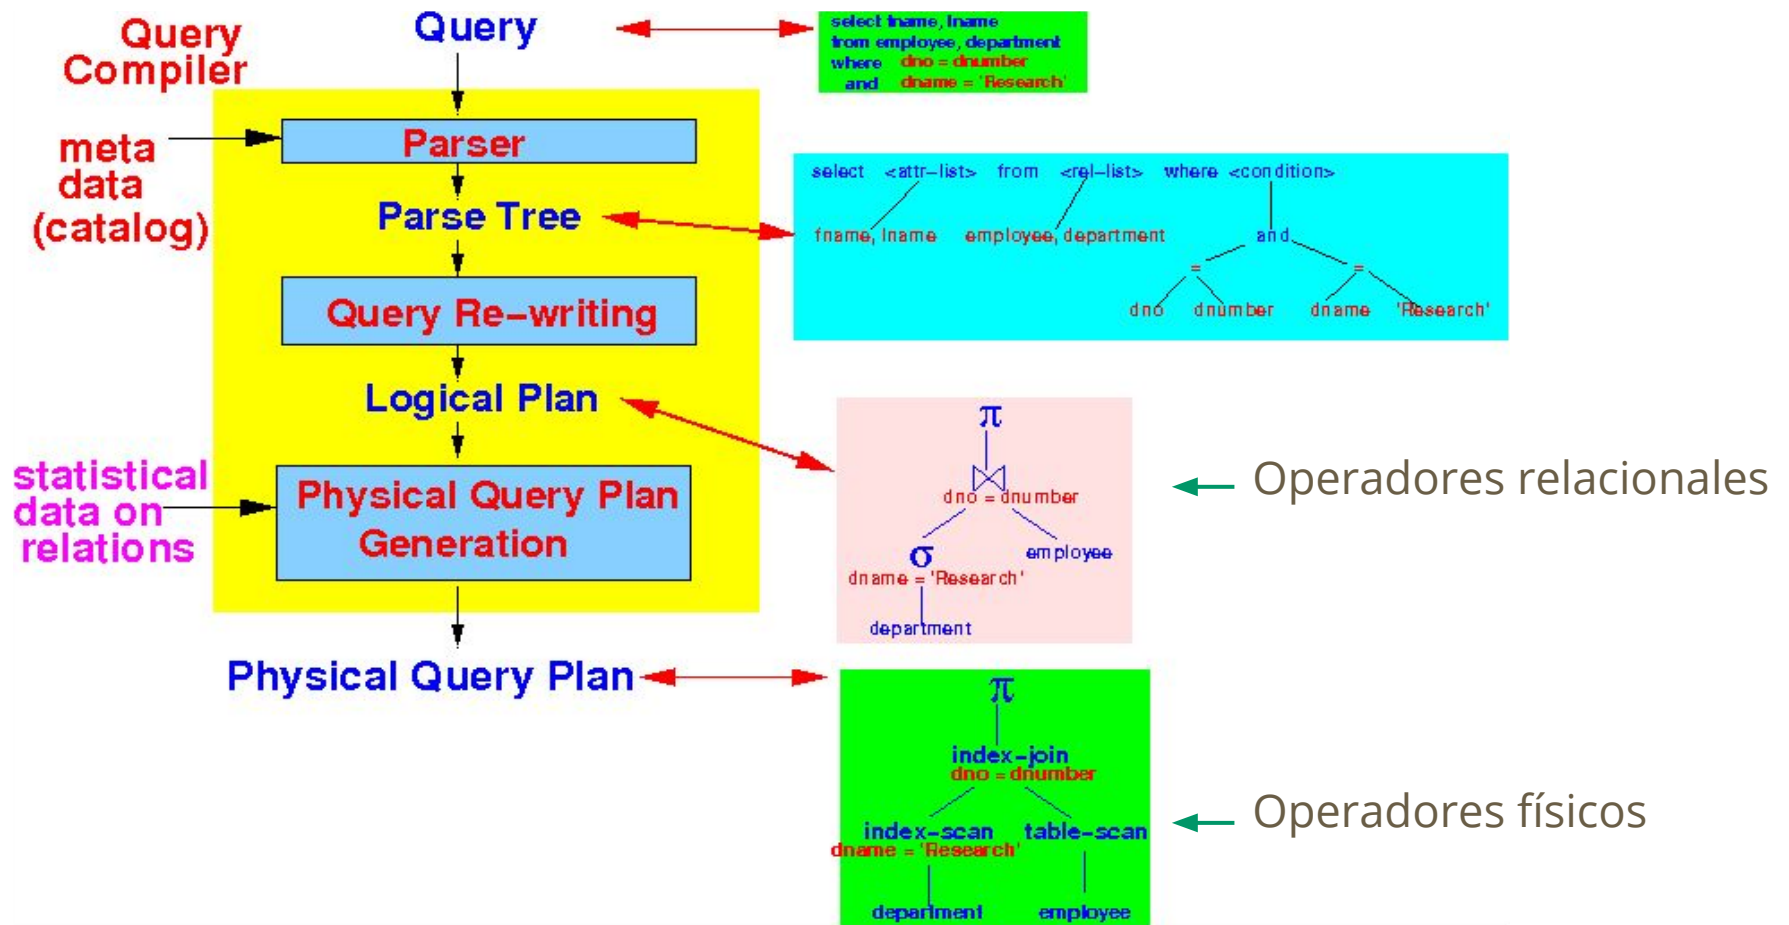
Symbol	Name of operator	Operation
$\text{table-scan}(R)$	Table scan	Reads blocks of the input relation r in - one block at a time.
$\text{index-scan}(R)$	Index scan	Reads blocks of the input relation r in using some index on the relation R
σ_{cond}	Selection	Selects the tuples that satisfies the condition cond
Π_{attrs}	Projection	Extracts the attributes in attrs from a tuple Note: result of Π is a bag (possible duplicates !!!)
\bowtie_{cond}	Theta Join	Joins 2 relations on the condition cond
\bowtie	Natural Join	Equi-joins 2 relations on the attributes than have the same name
\times	Product	Computes the (cartesian) product of 2 relations
δ	Duplicate Elimination	Removes the duplicate tuples from a bag of tuples (the result is a set)
γ_L	Grouping	Form gorups of tuples on common attribute values and apply the function L on each group
\cup_S	Set Union	Computes the set union of 2 sets (no duplicates)
\cup_B	Bag Union	Computes the bag union of 2 sets (allows duplicates)
\cap_S	Set Intersection	Computes the set intersection of 2 sets (no duplicates)
\cap_B	Bag Intersection	Computes the bag intersection of 2 sets (allows duplicates)
$-_S$	Set Difference	Computes the set difference of 2 sets (no duplicates)
$-_B$	Bag Difference	Computes the bag difference of 2 sets (allows duplicates)

Query Rewrite

Los planes de ejecución lógicos se expresan usando operadores relacionales

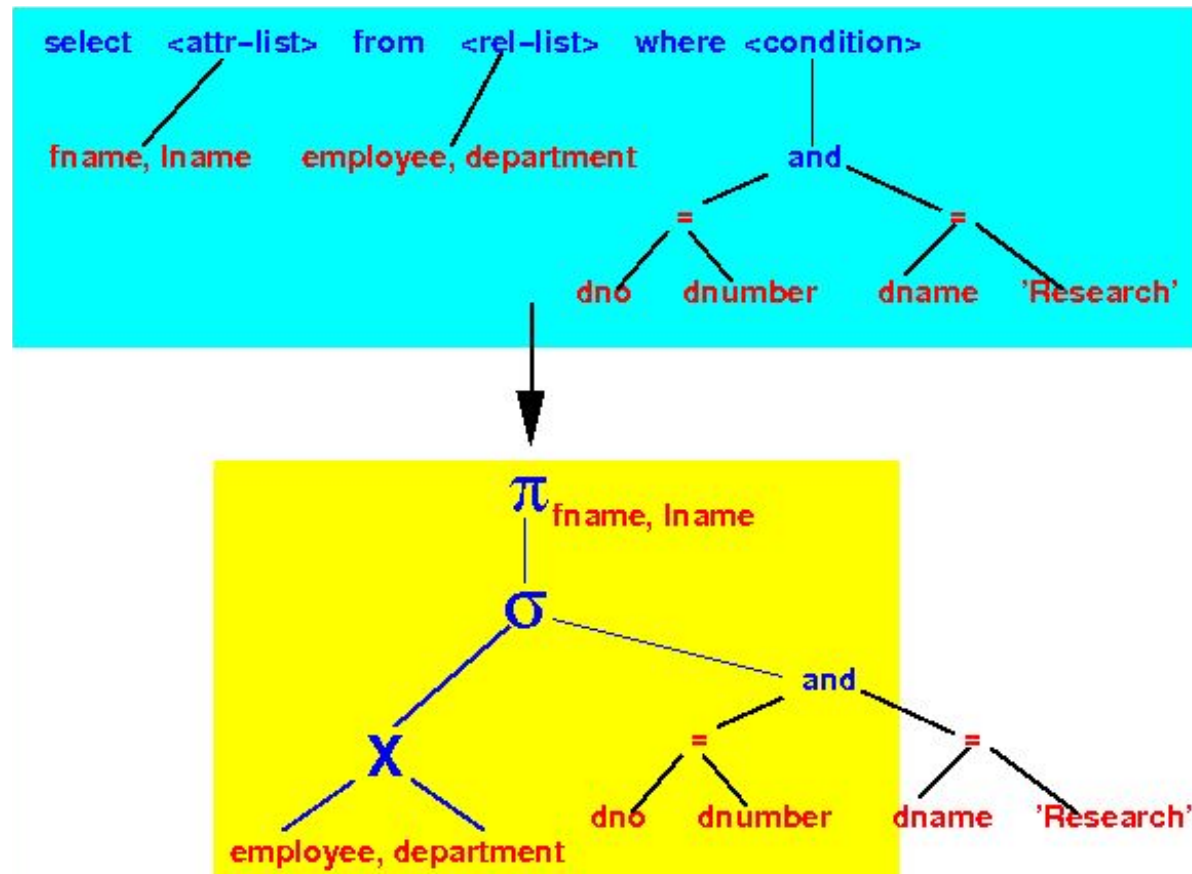
	Symbol	Name of operator	Operation
→	table-scan(R)	Table scan	Reads blocks of the input relation r in - one block at a time.
→	index-scan(R)	Index scan	Reads blocks of the input relation r in using some index on the relation R
→	σ_{cond}	Selection	Selects the tuples that satisfies the condition cond
→	Π_{attrs}	Projection	Extracts the attributes in attrs from a tuple Note: result of Π is a bag (possible duplicates !!!)
→	\bowtie_{cond}	Theta Join	Joins 2 relations on the condition cond
→	\bowtie	Natural Join	Equi-joins 2 relations on the attributes than have the same name
CROSS JOIN →	\times	Product	Computes the (cartesian) product of 2 relations
	δ	Duplicate Elimination	Removes the duplicate tuples from a bag of tuples (the result is a set)
GROUP BY →	γ_L	Grouping	Form gorups of tuples on common attribute values and apply the function L on each group
	\cup_S	Set Union	Computes the set union of 2 sets (no duplicates)
	\cup_B	Bag Union	Computes the bag union of 2 sets (allows duplicates)
	\cap_S	Set Intersection	Computes the set intersection of 2 sets (no duplicates)
	\cap_B	Bag Intersection	Computes the bag intersection of 2 sets (allows duplicates)
	$-_S$	Set Difference	Computes the set difference of 2 sets (no duplicates)
	$-_B$	Bag Difference	Computes the bag difference of 2 sets (allows duplicates)

Query Rewrite



Query Rewrite

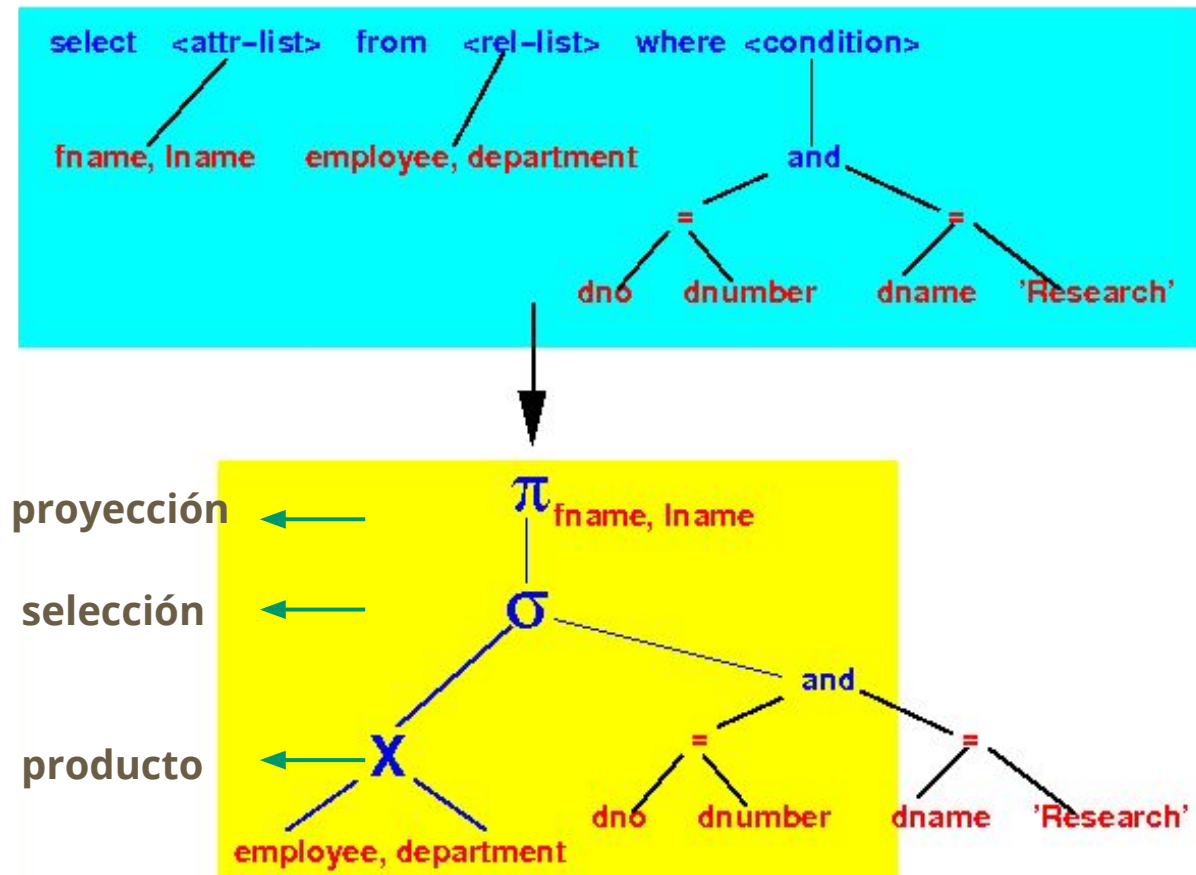
Parse-tree → plan lógico de la consulta



Partir desde el nivel más bajo hacia los niveles altos

Query Rewrite

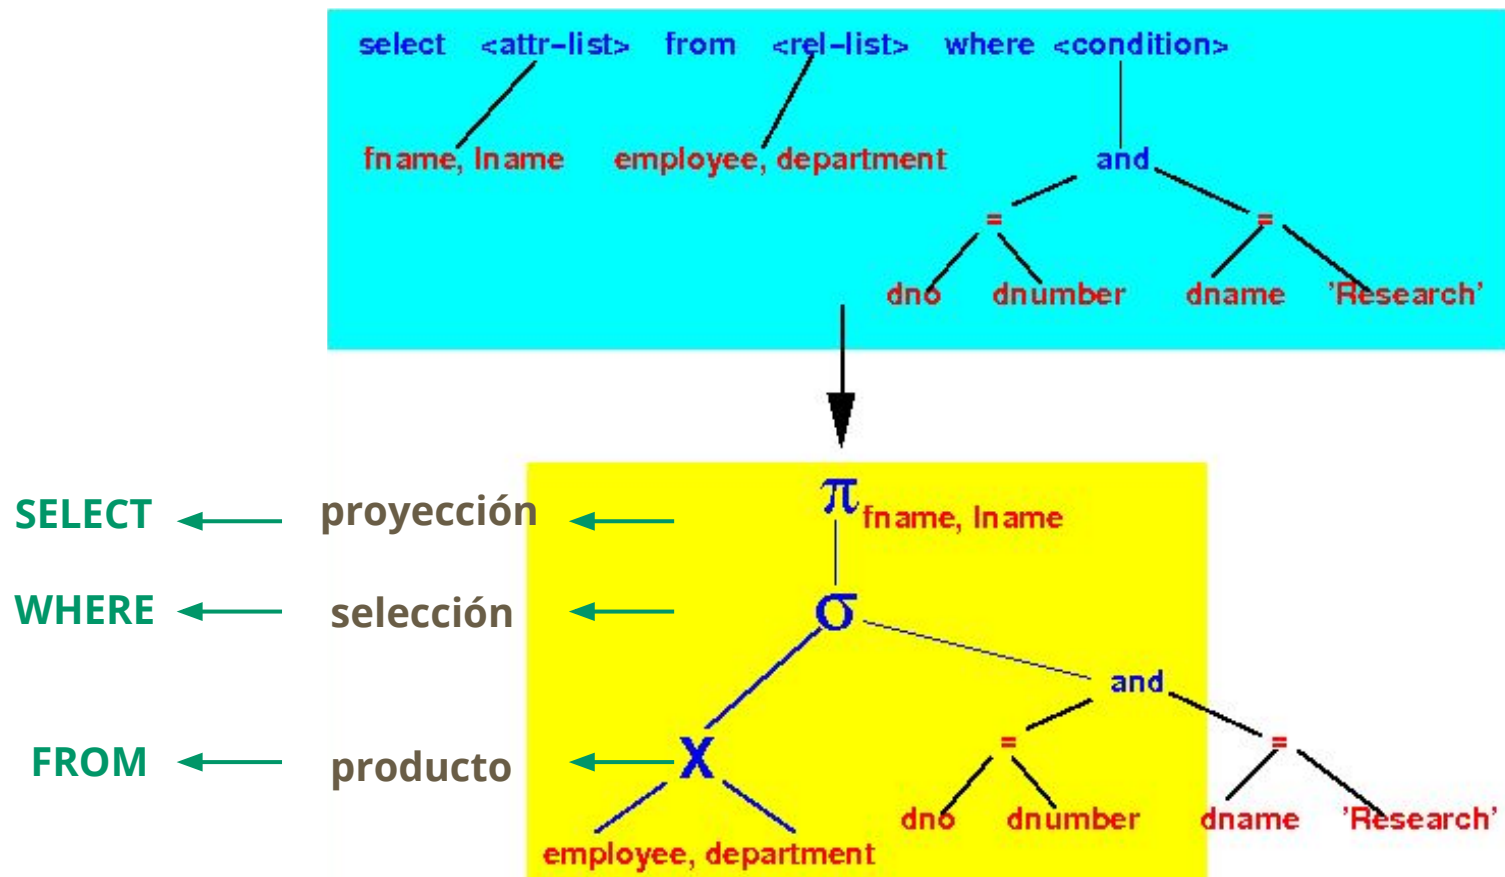
Parse-tree → plan lógico de la consulta



Partir desde el nivel más bajo hacia los niveles altos

Query Rewrite

Parse-tree → plan lógico de la consulta



Partir desde el nivel más bajo hacia los niveles altos

Query Planner

Plan lógico de la consulta → plan optimizado

Criterios de optimización:

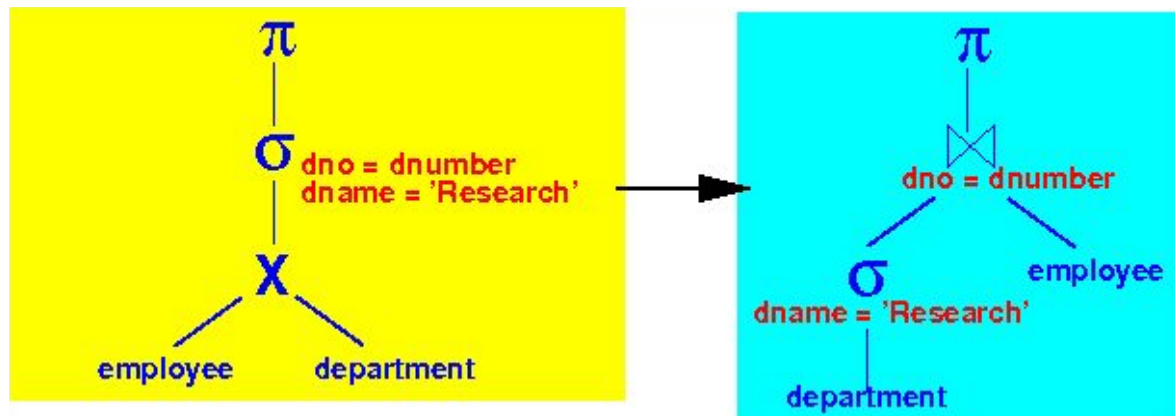
1. **Selection pushdown**: empujar la selección hacia las hojas
2. **Join reordering**: permutar las relaciones en un join

Query Planner

Plan lógico de la consulta → plan optimizado

Criterios de optimización:

1. **Selection pushdown**: empujar la selección hacia las hojas
2. **Join reordering**: permutar las relaciones en un join

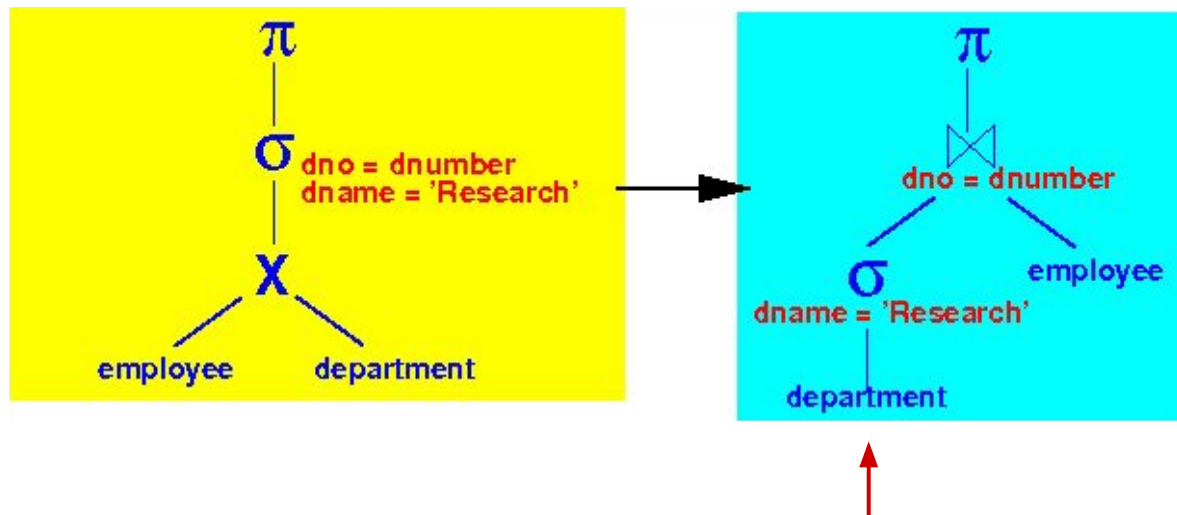


Query Planner

Plan lógico de la consulta → plan optimizado

Criterios de optimización:

1. **Selection pushdown**: empujar la selección hacia las hojas
2. **Join reordering**: permutar las relaciones en un join



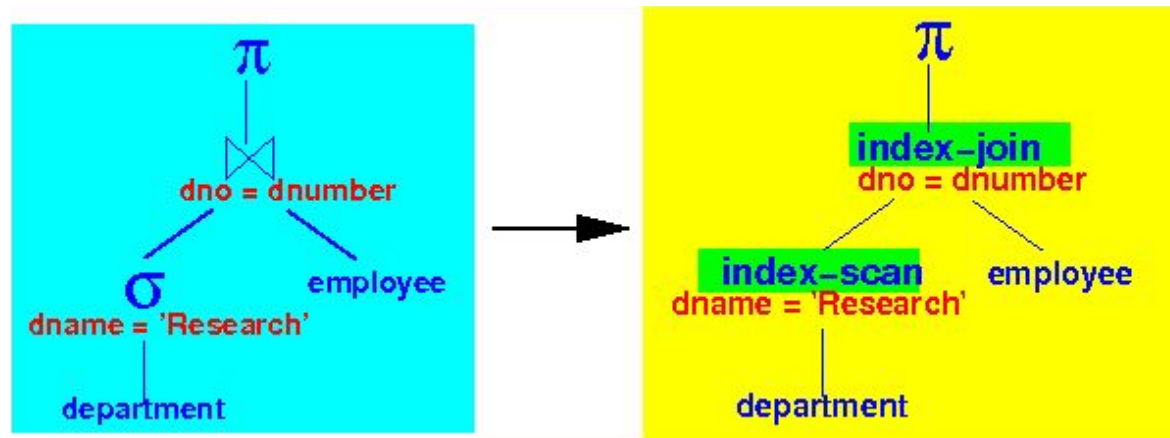
Aquí se hizo selection pushdown

Executor

Selecciona el mejor algoritmo para ejecutar el plan lógico optimizado

El **mejor algoritmo** depende de:

- Índices disponibles
- Memoria disponible para procesar la consulta



Executor

Selecciona el mejor algoritmo para ejecutar el plan lógico optimizado

El **mejor algoritmo** depende de:

- Índices disponibles
- Memoria disponible para procesar la consulta




Costo de una consulta

¿Qué variables afectan al
costo de ejecución?

Costo de una consulta

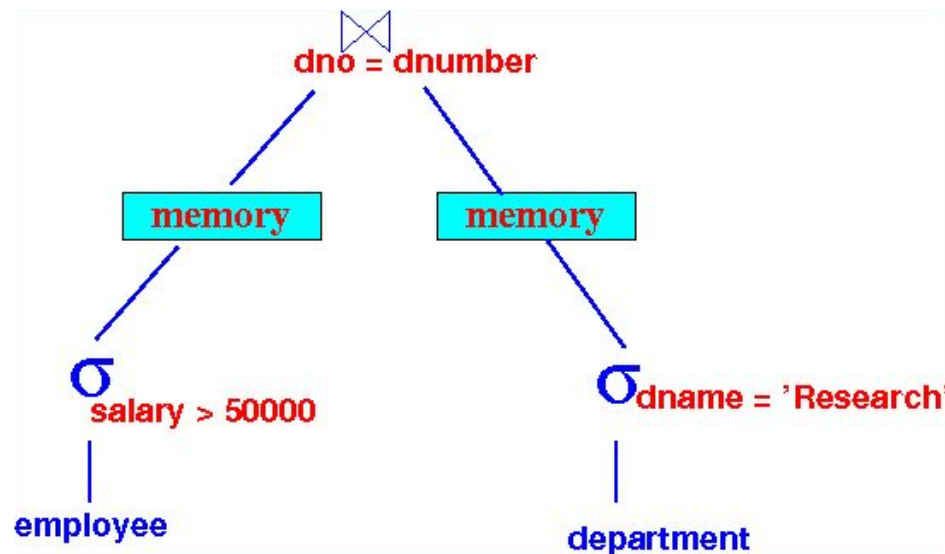
Una consulta requiere de *buffers* de memoria para poder computarse.



- La ejecución del plan físico requiere de varios *buffers* para alojar bloques de discos
- El costo de la consulta está determinado por el **número de bloques de disco accedidos**

Costo de una consulta


Asumimos que los resultados intermedios pueden serializarse en un pipeline, es decir, el operador que recibe los datos los lee desde un buffer:



- Los costos pueden contabilizarse como el **número de buffers** necesarios para transferir los resultados en el plan de la consulta
- Contabilizamos los costos en *buffers* de entrada y en memoria

Costo de operaciones básicas

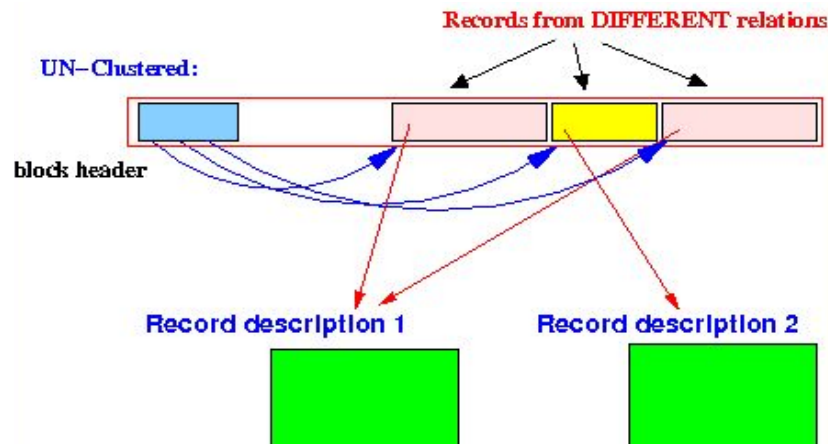
Archivos clusterizados



Operator	I/O cost	Explanation
Table-Scan	$B(R)$	<p>Read all blocks from the relation: $B(R)$ blocks.</p> <div> <p>Relation R:</p> </div>
Index-Scan	$\leq B(R)$	<p>Depends on number of values in the index is scanned</p> <p>Single value: I/O cost = $O(1)$ (a few blocks) $\ll B(R)$</p>

Costo de operaciones básicas

Archivos **no** clusterizados



Record header contains information of the relation to which the record belong

Operator	I/O cost	Explanation
Table-Scan	$\sim T(R)$	<p>Assuming that the tuples are sparsely stored in the file:</p> <div style="border: 1px solid black; padding: 10px; margin: 10px;"> <p>Relation R's tuples:</p> <div style="display: flex; justify-content: space-around; text-align: center;"> <div>1 block </div> <div>1 block </div> <div>1 block </div> <div>1 block </div> </div> <p style="text-align: center;">$B(R) = 8 \text{ blocks}$ $T(R) = 5 \text{ tuples}$</p> </div> <p>We will read 1 disk block per tuple. Total disk blocks read = $T(R)$</p>
Index-Scan	$\leq T(R)$	<p>Depends on number of values in the index is scanned</p> <p>Single value: I/O cost = $O(1)$ (a few blocks) $\ll B(R) \ll T(R)$</p>

Costo de operaciones básicas

Operador de selección

$\sigma_{\text{cond}} (R)$

Ejemplo:

$R = \{ ('John', 30000), ('Jane', 50000) \}$

$\sigma_{\text{salary} > 40000} (\text{employee}) = \{ ('Jane', 50000) \}$

Algoritmo de una pasada
(*one-pass*)

```
while ( R has more data blocks )
{
    read data blocks to fill available buffers;
    check condition for each tuple in the buffers;
    Move qualifying tuples to output
}
```


Costo de operaciones básicas

Operador de selección

$\sigma_{\text{cond}} (R)$

Ejemplo:


$R = \{ ('John', 30000), ('Jane', 50000) \}$

$\sigma_{\text{salary} > 40000} (\text{employee}) = \{ ('Jane', 50000) \}$

Algoritmo de una pasada
(*one-pass*)

```
while ( R has more data blocks )
{
    read data blocks to fill available buffers;
    check condition for each tuple in the buffers;
    Move qualifying tuples to output
}
```

Uso del *buffer*



Costos

- $B(R)$ --- if the relation R is *clustered*
- $T(R)$ --- if the relation R is *UNclustered*

Costo de operaciones básicas

Operador de proyección

```
 $\pi_{\text{attrs}} ( R )$ 
```

Ejemplo:

```
R = { ('John', 30000), ('Jane', 50000) }
```

```
 $\pi_{\text{fname}} ( \text{employee} ) = \{ ('John'), ('Jane') \}$ 
```

Algoritmo de una pasada
(*one-pass*)

```
while ( R has more data blocks )  
{  
    read data blocks to fill available buffers;  
    Extract the projection attrs for each tuple in buffers;  
    Move projected tuples to output  
}
```

Costo de operaciones básicas

Operador de proyección

Ejemplo:

$\pi_{\text{attrs}} (R)$


$R = \{ ('John', 30000), ('Jane', 50000) \}$

$\pi_{\text{fname}} (\text{employee}) = \{ ('John'), ('Jane') \}$

Algoritmo de una pasada
(*one-pass*)

```
while ( R has more data blocks )
{
    read data blocks to fill available buffers;
    Extract the projection attrs for each tuple in buffers;
    Move projected tuples to output
}
```

Uso del *buffer*



Costos

- $B(R)$ --- if the relation R is *clustered*
- $T(R)$ --- if the relation R is *UNclustered*

Costo de operaciones básicas

Operador de agrupamiento:

$$\gamma_L (R)$$

Ejemplo:

$$\begin{aligned} R &= (x, y) \\ R &= \{ (a, 2), (a, 3), (b, 1), (b, 3) \} \\ \gamma_{avg(y)} (R) &= \{ (a, 2.5), (b, 2.0) \} \end{aligned}$$

Costo de operaciones básicas

Algoritmo de una pasada (*one-pass*):

```
initialize a search structure H on grouping attributes of  $\gamma$ ;  
  
/* =====  
   Process the statistics for each group  
   ===== */  
while ( R has more data blocks )  
{  
    read 1 data block in buffer b;  
    for ( each tuple  $t \in b$  )  
    {  
        /* =====  
           We need a search structure H to implement the test  
            $t \in H$  efficiently !!!  
  
           We can use hash table or some bin. search tree  
           ===== */  
        if (  $t \in H$  )  
        {  
            Update the statistics for group(t);  
        }  
        else  
        {  
            insert t in H;  
  
            Initialize the statistics for group(t);  
        }  
    }  
}  
  
/* =====  
   Now we can output the aggregate function for each group  
   ===== */  
for ( each group  $\in H$  )  
{  
    Output group search key + statistics;  
}
```

- 1 *buffer* para leer las tuplas
- $M - 1$ *buffers* para almacenar el índice

Costo de operaciones básicas

Algoritmo de una pasada (*one-pass*):

```
initialize a search structure H on grouping attributes of  $\gamma$ ;  
  
/* =====  
   Process the statistics for each group  
   ===== */  
while ( R has more data blocks )  
{  
    read 1 data block in buffer b;  
    for ( each tuple  $t \in b$  )  
    {  
        /* =====  
           We need a search structure H to implement the test  
            $t \in H$  efficiently !!!  
  
           We can use hash table or some bin. search tree  
           ===== */  
        if (  $t \in H$  )  
        {  
            Update the statistics for group(t);  
        }  
        else  
        {  
            insert t in H;  
            Initialize the statistics for group(t);  
        }  
    }  
}  
  
/* =====  
   Now we can output the aggregate function for each group  
   ===== */  
for ( each group  $\in H$  )  
{  
    Output group search key + statistics;  
}
```

- 1 *buffer* para leer las tuplas
- $M - 1$ *buffers* para almacenar el índice

Costo:

- $B(R)$ --- if the relation R is *clustered*
- $T(R)$ --- if the relation R is *UNclustered*

Memoria:

$$M = O(V(R, [A_1, A_2, \dots, A_k]))$$

donde:

- $V(R, [A_1, A_2, \dots, A_k]) =$ #valores que $[A_1, \dots, A_k]$ puede tomar

Consultas?

Recuerden!

- Marcelo Mendoza: mmendoza@inf.utfsm.cl
- Margarita Bugueño: margarita.bugueno@usm.cl

Bases de Datos

— Ejecución de consultas —